

A post-CM infrastructure delivery pipeline

... or why I'm here to learn more

@beddari at @cfgmgtcamp

Problem statement

We were using CM tooling but
NOT WINNING

What we had built with love -

- √ automated tests
- √ JEOS + baseline + role
- √ monitoring across envs

– was a total failure!

- √ non-manageable rebuild times
- √ envs were starting to “leak”
- √ upgrades were high risk

“Our systems are ...
eventually repeatable”

Darn it, test that small change in prod

People told us
“CLEARLY
you are doing something wrong”



dock rocker docker rocker docker docker
docker rocker d docker rocker
docker doc ker docke

Solution:

We stopped doing
configuration management

input / change / output

input

Inputs are typically managed artifacts

Repos, packages, images, containers

change

Feed input to Packer which in turn runs a
builder that applies *change*,
producing output

output

A versioned artifact, suitable for consumption
or further processing

Repos, packages, images, containers

Abstraction

is key

A input-change-output chain is a *project*

•

A project is *versioned in git*

•

Artifacts are *testable*

Your new job is ...

describing state to produce artifacts
and keeping that state from drifting

We didn't throw it all out

A *system* consists of 1+N roles,
a *role* consists of 1+N profiles.

First lesson learned

Think about managing state,
where do you want it?

<https://github.com/Nubisproject/nubis-docs/blob/master/MANIFEST0.md>

Abstraction is future proofing

<http://jerakia.io>

“Decoupled from any particular configuration management system”

Does it work?

YES

Packer with masterless Puppet
Terraform and Ansible to deploy and replace
Masterless Puppet to audit and correct drift

Single YAML configuration data store

yum upgrade considered harmful